

Lecture 11

Common Optimization Algorithms

STAT 453: Deep Learning, Spring 2020

Sebastian Raschka

<http://stat.wisc.edu/~sraschka/teaching/stat453-ss2020/>

Slides:

<https://github.com/rasbt/stat453-deep-learning-ss20/tree/master/L11-optim>

Course overview

Part 1: Introduction

- [L01: Course overview, introduction to deep learning](#)
- [L02: The brief history of deep learning](#)
- [L03: Single-layer neural networks: The perceptron algorithm](#)

Part 2: Mathematical and computational foundations

- [L04: Linear algebra and calculus for deep learning](#)
- [L05: Parameter optimization with gradient descent](#)
- [L06: Automatic differentiation & PyTorch](#)
- [L06.5: Cloud Computing \[Recording \]](#)

Part 3: Introduction to neural networks

- [L07: Multinomial logistic regression](#)
- [L08: Multilayer perceptrons](#)
- [L09: Regularization \[Recording \]](#)
- [L10: Input normalization and weight initialization \[Recording Part 1/2 \] \[Recording Part 2/2 \]](#)
- Learning rates and advanced optimization algorithms

Part 4: Deep learning for computer vision and language modeling

- Introduction to convolutional neural networks 1
- Introduction to convolutional neural networks 2
- Introduction to recurrent neural networks 1
- Introduction to recurrent neural networks 2

Part 5: Deep generative models

- Autoencoders
- Autoregressive models
- Variational autoencoders
- Normalizing Flow Models
- Generative adversarial networks 1
- Generative adversarial networks 2
- Evaluating generative models

Overview: Additional Tricks for Neural Network Training (Part 2/2)

Part 1 (Last Lecture, L10)

- Input Normalization (BatchNorm, InstanceNorm, GroupNorm, LayerNorm)
- Weight Initialization (Xavier Glorot, Kaiming He)

Part 2 (this lecture)

- Learning Rate Decay
- Momentum Learning
- Adaptive Learning

Overview: Additional Tricks for Neural Network Training (Part 2/2)

Part 1 (Last Lecture, L10)

- Input Normalization (BatchNorm, InstanceNorm, GroupNorm, LayerNorm)
- Weight Initialization (Xavier Glorot, Kaiming He)

Part 2 (this lecture)

- Learning Rate Decay
 - Momentum Learning
 - Adaptive Learning
- (Modifications of the 1st order SGD optimization algorithm; 2nd order methods are rarely used in DL)

Improving Stochastic Gradient Descent

Lecture Overview

1. Applying "tricks" to HW3
2. Learning rate decay
3. Training with "momentum"
4. Adaptive learning rates
5. Using optimization algorithms in PyTorch

Experiments with HW03 code based on "tricks" from last lecture

```
class MLP(torch.nn.Module):



    def __init__(self, num_features, num_hidden_1,
                 num_hidden_2, num_hidden_3, num_classes):
        super(MLP, self).__init__()
        self.num_classes = num_classes
        self.linear_1 = torch.nn.Linear(num_features, num_hidden_1)
        self.linear_2 = torch.nn.Linear(num_hidden_1, num_hidden_2)
        self.linear_3 = torch.nn.Linear(num_hidden_2, num_hidden_3)
        self.linear_out = torch.nn.Linear(num_hidden_3, num_classes)

    def forward(self, x):
        out = self.linear_1(x)
        out = torch.tanh(out)
        out = self.linear_2(out)
        out = torch.tanh(out)
        out = self.linear_3(out)
        out = torch.tanh(out)
        logits = self.linear_out(out)
        probas = F.softmax(logits, dim=1)
        return logits, probas

model = MLP(num_features=28*28,
            num_hidden_1=75,
            num_hidden_2=50,
            num_hidden_3=25,
            num_classes=10)
```







Experiments with HW03 code based on "tricks" from last lecture

01_hw3-tanh-uniform.ipynb







<https://github.com/rasbt/stat453-deep-learning-ss20/tree/master/L11-optim/code/HW03-experiments>

Experiments with HW03 code based on "tricks" from last lecture



Experiments with HW03 code based on "tricks" from last lecture




Improving Stochastic Gradient Descent

Lecture Overview


1. Applying "tricks" to HW3
- 2. Learning rate decay**
3. Training with "momentum"
4. Adaptive learning rates
5. Using optimization algorithms in PyTorch

Minibatch Learning Recap



- Minibatch learning is a form of stochastic gradient descent
- Each minibatch can be considered a sample drawn from the training set (where the training set is in turn a sample drawn from the population)
- Hence, the gradient is noisier
- A noisy gradient can be
 - ◆ good: chance to escape local minima
 - ◆ bad: can lead to extensive oscillation
- Main advantage: Convergence speed, because it offers opportunities for parallelism (do you recall what these are?)


Minibatch Learning Recap



- Minibatch learning is a form of stochastic gradient descent
- Each minibatch can be considered a sample drawn from the training set (where the training set is in turn a sample drawn from the population)
- Hence, the gradient is noisier
- A noisy gradient can be
 - ◆ good: chance to escape local minima
 - ◆ bad: can lead to extensive oscillation
- Main advantage: Convergence speed, because it offers two opportunities for parallelism (do you recall what these are?)
- Note that second order methods that take e.g., gradient curvature into account usually don't work so well in practice and are not often used/recommended in DL

Nice Library & Visualization Tool

<https://vis.esmallen.org>



Practical Tip for Minibatch Use

- Reasonable minibatch sizes are usually: 32, 64, 128, 256, 512, 1024 (in the last lecture, we discussed why powers of 2 are a common convention)
- Usually, you can choose a batch size that is as large as your GPU memory allows (matrix-multiplication and the size of fully-connected layers are usually the bottleneck)
- Practical tip: usually, it is a good idea to also make the batch size proportional to the number of classes in the dataset





Figure 1: Distribution of *Iris* flower classes upon random subsampling into training and test sets.

Raschka, S. (2018). Model evaluation, model selection, and algorithm selection in machine learning. *arXiv preprint arXiv:1811.12808*.


Learning Rate Decay

- Batch effects -- minibatches are samples of the training set, hence minibatch loss and gradients are approximations
- Hence, we usually get oscillations
- To dampen oscillations towards the end of the training, we can decay the learning rate



Learning Rate Decay

- Batch effects -- minibatches are samples of the training set, hence minibatch loss and gradients are approximations
- Hence, we usually get oscillations
- To dampen oscillations towards the end of the training, we can decay the learning rate



Danger of learning rate is to decrease the learning rate too early

Practical tip: try to train the model without learning rate decay first, then add it later

You can also use the validation performance (e.g., accuracy) to judge whether lr decay is useful (as opposed to using the training loss)


Learning Rate Decay

Most common variants for learning rate decay:

1) Exponential Decay:

$$\eta_t := \eta_0 \cdot e^{-k \cdot t}$$

where k is the decay rate



Learning Rate Decay


Most common variants for learning rate decay:

2) Halving the learning rate:

$$\eta_t := \eta_{t-1}/2$$

3) Inverse decay:

$$\eta_t := \frac{\eta_0}{1 + k \cdot t}$$



Learning Rate Decay

There are many, many more

E.g., Cyclical Learning Rate

Smith, Leslie N. “[Cyclical learning rates for training neural networks](#).” Applications of Computer Vision (WACV), 2017 IEEE Winter Conference on. IEEE, 2017.




Figure 2. Triangular learning rate policy. The blue lines represent learning rate values changing between bounds. The input parameter *stepsize* is the number of iterations in half a cycle.

(which, I found, didn't work well at all in practice, unfortunately -- at least in my case)

Learning Rate Decay in PyTorch

Option 1. Just call your own function at the end of each epoch:

```
def adjust_learning_rate(optimizer, epoch, initial_lr, decay_rate):
    """Exponential decay every 10 epochs"""
    if not epoch % 10:
        lr = initial_lr * torch.exp(-decay_rate*epoch)
        for param_group in optimizer.param_groups:
            param_group['lr'] = lr
```

Learning Rate Decay in PyTorch

Option 2. Use one of the built-in tools in PyTorch: (many more available)
(Here, the most generic version.)

CLASS `torch.optim.lr_scheduler.LambdaLR(optimizer, lr_lambda, last_epoch=-1)`

[SOURCE]

Sets the learning rate of each parameter group to the initial lr times a given function. When `last_epoch=-1`, sets initial lr as lr.

- Parameters:**
- **optimizer** (*Optimizer*) – Wrapped optimizer.
 - **lr_lambda** (*function or list*) – A function which computes a multiplicative factor given an integer parameter epoch, or a list of such functions, one for each group in `optimizer.param_groups`.
 - **last_epoch** (*int*) – The index of last epoch. Default: -1.

Example

```
>>> # Assuming optimizer has two groups.  
>>> lambda1 = lambda epoch: epoch // 30  
>>> lambda2 = lambda epoch: 0.95 ** epoch  
>>> scheduler = LambdaLR(optimizer, lr_lambda=[lambda1, lambda2])  
>>> for epoch in range(100):  
>>>     scheduler.step()  
>>>     train(...)  
>>>     validate(...)
```

Source: <https://pytorch.org/docs/stable/optim.html>

Learning Rate Decay in PyTorch

Example, part 1/2

```
#####
### Model Initialization
#####

torch.manual_seed(RANDOM_SEED)
model = MLP(num_features=28*28,
            num_hidden=100,
            num_classes=10)

model = model.to(DEVICE)

optimizer = torch.optim.SGD(model.parameters(), lr=0.1)

#####
### LEARNING RATE SCHEDULER
#####

scheduler = torch.optim.lr_scheduler.ExponentialLR(optimizer,
                                                    gamma=0.1,
                                                    last_epoch=-1)

...
```

Learning Rate Decay in PyTorch

```
for epoch in range(5):
    model.train()
    for batch_idx, (features, targets) in enumerate(train_loader):

        features = features.view(-1, 28*28).to(DEVICE)
        targets = targets.to(DEVICE)

        ### FORWARD AND BACK PROP
        logits, probas = model(features)

        #cost = F.nll_loss(torch.log(probas), targets)
        cost = F.cross_entropy(logits, targets)
        optimizer.zero_grad()

        cost.backward()
        minibatch_cost.append(cost)
        ### UPDATE MODEL PARAMETERS

        optimizer.step()

        ### LOGGING
        if not batch_idx % 50:
            print ('Epoch: %03d/%03d | Batch %03d/%03d | Cost: %.4f'
                  %(epoch+1, NUM_EPOCHS, batch_idx,
                    len(train_loader), cost))

#####
### Update Learning Rate
scheduler.step() # don't have to do it every epoch!
#####

model.eval()
```

Example, part 2/2

Saving Models in PyTorch

Save Model

```
model.to(torch.device('cpu'))
torch.save(model.state_dict(), './my_model_2epochs.pt')
torch.save(optimizer.state_dict(), './my_optimizer_2epochs.pt')
torch.save(scheduler.state_dict(), './my_scheduler_2epochs.pt')
```

Load Model

```
model = MLP(num_features=28*28,
            num_hidden=100,
            num_classes=10)

model.load_state_dict(torch.load('./my_model_2epochs.pt'))
model = model.to(DEVICE)

# for this particular optimizer not necessary, as it doesn't have a state
# but good practice, so you don't forget it when using other optimizers
# later
optimizer = torch.optim.SGD(model.parameters(), lr=0.1)
optimizer.load_state_dict(torch.load('./my_optimizer_2epochs.pt'))

scheduler = torch.optim.lr_scheduler.ExponentialLR(optimizer,
                                                    gamma=0.1,
                                                    last_epoch=-1)
scheduler.load_state_dict(torch.load('./my_scheduler_2epochs.pt'))

model.train()
```

Learning rate schedulers have the advantage that we can also simply save their state for reuse (e.g., saving and continuing training later)

https://github.com/rasbt/stat479-deep-learning-ss19/tree/master/L12_optim/lr_scheduler_and_saving_models.ipynb

Relationship between Learning Rate and Batch Size

DON'T DECAY THE LEARNING RATE,
INCREASE THE BATCH SIZE

Samuel L. Smith*, Pieter-Jan Kindermans*, Chris Ying & Quoc V. Le

Google Brain


{slsmith, pikinder, chrisying, qvl}@google.com

ABSTRACT


It is common practice to decay the learning rate. Here we show one can usually obtain the same learning curve on both training and test sets by instead increasing the batch size during training. This procedure is successful for stochastic gradient descent (SGD), SGD with momentum, Nesterov momentum, and Adam. It reaches equivalent test accuracies after the same number of training epochs, but with fewer parameter updates, leading to greater parallelism and shorter training times. We can further reduce the number of parameter updates by increasing the learning rate ϵ and scaling the batch size $B \propto \epsilon$. Finally, one can increase the momentum coefficient m and scale $B \propto 1/(1 - m)$, although this tends to slightly reduce the test accuracy. Crucially, our techniques allow us to repurpose existing training schedules for large batch training with no hyper-parameter tuning. We train ResNet-50 on ImageNet to 76.1% validation accuracy in under 30 minutes.

Smith, S. L., Kindermans, P. J., Ying, C., & Le, Q. V. (2017). Don't decay the learning rate, increase the batch size. *arXiv preprint arXiv:1711.00489*.

Relationship between Learning Rate and Batch Size



(a)



(b)

Figure 6: Inception-ResNet-V2 on ImageNet. Increasing the batch size during training achieves similar results to decaying the learning rate, but it reduces the number of parameter updates from just over 14000 to below 6000. We run each experiment twice to illustrate the variance.

Smith, S. L., Kindermans, P. J., Ying, C., & Le, Q. V. (2017). Don't decay the learning rate, increase the batch size. *arXiv preprint arXiv:1711.00489*.

Improving Stochastic Gradient Descent

Lecture Overview

1. Applying "tricks" to HW3
2. Learning rate decay
- 3. Training with "momentum"**
4. Adaptive learning rates
5. Using optimization algorithms in PyTorch

Training with "Momentum"

Momentum

From Wikipedia, the free encyclopedia

This article is about linear momentum. It is not to be confused with angular momentum.

This article is about momentum in physics. For other uses, see Momentum (disambiguation).


In [Newtonian mechanics](#), **linear momentum**, **translational momentum**, or simply **momentum** (pl. momenta) is the product of the [mass](#) and [velocity](#) of an object. It is a [vector](#) quantity, possessing a magnitude and a direction in three-dimensional space. If m is an object's mass and \mathbf{v} is the velocity (also a vector), then the momentum is

Source: <https://en.wikipedia.org/wiki/Momentum>


- Concept: In momentum learning, we try to accelerate convergence by dampening oscillations using "velocity" (the speed of the "movement" from previous updates)

Training with "Momentum"

- Concept: In momentum learning, we try to accelerate convergence by dampening oscillations using "velocity" (the speed of the "movement" from previous updates)




Without momentum




With momentum

Training with "Momentum"



Without momentum



With momentum

Key take-away:


Not only move in the (opposite) direction of the gradient, but also move in the "averaged" direction of the last few updates

Training with "Momentum"

Key take-away:

Not only move in the (opposite) direction of the gradient, but also move in the "averaged" direction of the last few updates

Helps with dampening oscillations, but also helps with escaping local minima traps



Training with "Momentum"

Often referred to as "velocity" v

$$\Delta w_{i,j}(t) := \alpha \cdot \Delta w_{i,j}(t - 1) + \eta \cdot \frac{\partial \mathcal{L}}{\partial w_{i,j}}(t)$$

Usually, we choose a momentum rate between 0.9 and 0.999; you can think of it as a "friction" or "dampening" parameter


"velocity" from the previous iteration

Regular partial derivative/gradient multiplied by learning rate at current time step t

Weight update using the velocity vector:

$$w_{i,j}(t + 1) := w_{i,j}(t) - \Delta w_{i,j}(t)$$


Qian, N. (1999). On the momentum term in gradient descent learning algorithms. *Neural Networks : The Official Journal of the International Neural Network Society*, 12(1), 145–151. [http://doi.org/10.1016/S0893-6080\(98\)00116-6](http://doi.org/10.1016/S0893-6080(98)00116-6)



Step-size $\alpha = 0.02$



We often think of Momentum as a means of dampening oscillations and speeding up the iterations, leading to faster convergence. But it has other interesting behavior. It allows a larger range of step-sizes to be used, and creates its own oscillations. What is going on?



Step-size $\alpha = 0.02$



We often think of Momentum as a means of dampening oscillations and speeding up the iterations, leading to faster convergence. But it has other interesting behavior. It allows a larger range of step-sizes to be used, and creates its own oscillations. What is going on?

Source: <https://distill.pub/2017/momentum/>

CLASS `torch.optim.SGD(params, lr=<required parameter>, momentum=0, dampening=0, weight_decay=0, nesterov=False)`

[SOURCE]

Implements stochastic gradient descent (optionally with momentum).

Nesterov momentum is based on the formula from [On the importance of initialization and momentum in deep learning](#).

Parameters:

- **params** (*iterable*) – iterable of parameters to optimize or dicts defining parameter groups
- **lr** (*float*) – learning rate
- **momentum** (*float, optional*) – momentum factor (default: 0)
- **weight_decay** (*float, optional*) – weight decay (L2 penalty) (default: 0)
- **dampening** (*float, optional*) – dampening for momentum (default: 0)
- **nesterov** (*bool, optional*) – enables Nesterov momentum (default: False)

Example

Source: <https://pytorch.org/docs/stable/optim.html>

```
CLASS torch.optim.SGD(params, lr=<required parameter>, momentum=0, dampening=0, weight_decay=0,  
nesterov=False)
```

[SOURCE]

Implements stochastic gradient descent (optionally with momentum).

Nesterov momentum is based on the formula from [On the importance of initialization and momentum in deep learning](#).

- Parameters:**
- **params** (*iterable*) – iterable of parameters to optimize or dicts defining parameter groups
 - **lr** (*float*) – learning rate
 - **momentum** (*float, optional*) – momentum factor (default: 0)
 - **weight_decay** (*float, optional*) – weight decay (L2 penalty) (default: 0)
 - **dampening** (*float, optional*) – dampening for momentum (default: 0)
 - **nesterov** (*bool, optional*) – enables Nesterov momentum (default: False)

Example

Note that the optional "dampening" term is used as follows:

```
v = momentum * v + (1-dampening) * gradientW  
w = w - lr * v
```

Also note that in PyTorch, the learning rate is also applied to the momentum terms, instead of the original definition, which would be

```
v = momentum * v + (1-dampening) * lr * gradientW  
w = w - v
```

A Better Momentum Method: Nesterov Accelerated Gradient

Similar to momentum learning, but performs a correction after the update (based on where the loss, w.r.t. the weight parameters, is approx. going to be after the update)

Before:

$$\begin{aligned}\Delta \mathbf{w}_t &:= \alpha \cdot \Delta \mathbf{w}_{t-1} + \eta \cdot \nabla_{\mathbf{w}} \mathcal{L}(\mathbf{w}_t) \\ \mathbf{w}_{t+1} &:= \mathbf{w}_t - \Delta \mathbf{w}_t\end{aligned}$$

Nesterov:

$$\begin{aligned}\Delta \mathbf{w}_t &:= \alpha \cdot \Delta \mathbf{w}_{t-1} + \eta \cdot \nabla_{\mathbf{w}} \mathcal{L}(\mathbf{w}_t - \alpha \cdot \Delta \mathbf{w}_{t-1}) \\ \mathbf{w}_{t+1} &:= \mathbf{w}_t - \Delta \mathbf{w}_t\end{aligned}$$

Nesterov, Y. (1983). A method for unconstrained convex minimization problem with the rate of convergence $\mathcal{o}(1/k^2)$.
Doklady ANSSSR (translated as Soviet.Math.Docl.), vol. 269, pp. 543– 547.

Sutskever, I., Martens, J., Dahl, G. E., & Hinton, G. E. (2013). On the importance of initialization and momentum in deep learning. *ICML* (3), 28(1139-1147), 5.

A Better Momentum Method: Nesterov Accelerated Gradient




Figure 1. (Top) Classical Momentum (Bottom) Nesterov Accelerated Gradient

Sutskever, I., Martens, J., Dahl, G. E., & Hinton, G. E. (2013). On the importance of initialization and momentum in deep learning. *ICML* (3), 28(1139-1147), 5.

Improving Stochastic Gradient Descent

Lecture Overview

1. Applying "tricks" to HW3
2. Learning rate decay
3. Training with "momentum"
- 4. Adaptive learning rates**
5. Using optimization algorithms in PyTorch

Adaptive Learning Rates

There are many different flavors of adapting the learning rate
(bit out of scope for this course to review them all)

Key take-aways:

- decrease learning if the gradient changes its direction
- increase learning if the gradient stays consistent

Adaptive Learning Rates

Key take-aways:

- decrease learning if the gradient changes its direction
- increase learning if the gradient stays consistent

Step 1: Define a local gain (g) for each weight (initialized with $g=1$)

$$\Delta w_{i,j} := \eta \cdot g_{i,j} \cdot \frac{\partial \mathcal{L}}{\partial w_{i,j}}$$

Adaptive Learning Rates

Step 1: Define a local gain (g) for each weight (initialized with $g=1$)

$$\Delta w_{i,j} := \eta \cdot g_{i,j} \cdot \frac{\partial \mathcal{L}}{\partial w_{i,j}}$$

Step 2:

If gradient is consistent

$$g_{i,j}(t) := g_{i,j}(t - 1) + \beta$$

else

$$g_{i,j}(t) := g_{i,j}(t - 1) \cdot (1 - \beta)$$

Note that

multiplying by a factor has a larger impact if gains are large, compared to adding a term

(dampening effect if updates oscillate in the wrong direction)

Adaptive Learning Rate via RMSProp

- Unpublished algorithm by Geoff Hinton (but very popular) based on Rprop [1]
- Very similar to another concept called AdaDelta
- Concept: divide learning rate by exponentially decreasing moving average of the squared gradients
- This takes into account that gradients can vary widely in magnitude
- Here, RMS stands for "Root Mean Squared"
- Also, damps oscillations like momentum (but in practice, works a bit better)

[1] Igel, Christian, and Michael Hüsker. "Improving the Rprop learning algorithm." *Proceedings of the Second International ICSC Symposium on Neural Computation (NC 2000)*. Vol. 2000. ICSC Academic Press, 2000.

Adaptive Learning Rate via RMSProp

$$MeanSquare(w_{i,j}, t) := \beta \cdot MeanSquare(w_{i,j}, t - 1) + (1 - \beta) \left(\frac{\partial \mathcal{L}}{w_{i,j}(t)} \right)^2$$

moving average of the squared gradient for each weight

$$w_{i,j}(t) := w_{i,j}(t) - \eta \cdot \frac{\partial \mathcal{L}}{\partial w_{i,j}(t)} / \left(\sqrt{MeanSquare(w_{i,j}, t)} + \epsilon \right)$$

where beta is typically between 0.9 and 0.999

small epsilon term to
avoid division by zero

Adaptive Learning Rate via ADAM

- ADAM (Adaptive Moment Estimation) is probably the most widely used optimization algorithm in DL as of today
- It is a combination of the momentum method and RMSProp

Momentum-like term:

$$\Delta w_{i,j}(t) := \alpha \cdot \Delta w_{i,j}(t-1) + \eta \cdot \frac{\partial \mathcal{L}}{\partial w_{i,j}}(t)$$

m_{t-1}

original momentum term

$$m_t := \alpha \cdot m_{t-1} + (1 - \alpha) \cdot \frac{\partial \mathcal{L}}{\partial w_{i,j}}(t)$$

m_t

Kingma, D. P., & Ba, J. (2014). Adam: A method for stochastic optimization. *arXiv preprint arXiv:1412.6980*.

Adaptive Learning Rate via ADAM

Momentum-like term:

$$m_t := \alpha \cdot m_{t-1} + (1 - \alpha) \cdot \frac{\partial \mathcal{L}}{\partial w_{i,j}}(t)$$

RMSProp term:

$$r := \beta \cdot \text{MeanSquare}(w_{i,j}, t-1) + (1 - \beta) \left(\frac{\partial \mathcal{L}}{\partial w_{i,j}}(t) \right)^2$$

ADAM update:

$$w_{i,j} := w_{i,j} - \eta \frac{m_t}{\sqrt{r} + \epsilon}$$

Kingma, D. P., & Ba, J. (2014). Adam: A method for stochastic optimization. *arXiv preprint arXiv:1412.6980*.

Adaptive Learning Rate via ADAM

Algorithm 1: *Adam*, our proposed algorithm for stochastic optimization. See section 2 for details, and for a slightly more efficient (but less clear) order of computation. g_t^2 indicates the elementwise square $g_t \odot g_t$. Good default settings for the tested machine learning problems are $\alpha = 0.001$, $\beta_1 = 0.9$, $\beta_2 = 0.999$ and $\epsilon = 10^{-8}$. All operations on vectors are element-wise. With β_1^t and β_2^t we denote β_1 and β_2 to the power t .

Require: α : Step size

Require: $\beta_1, \beta_2 \in [0, 1]$: Exponential decay rates for the moment estimates

Require: $f(\theta)$: Stochastic objective function with parameters θ

Require: θ_0 : Initial parameter vector

$m_0 \leftarrow 0$ (Initialize 1st moment vector)

$v_0 \leftarrow 0$ (Initialize 2nd moment vector)

$t \leftarrow 0$ (Initialize timestep)

while θ_t not converged **do**

$t \leftarrow t + 1$

$g_t \leftarrow \nabla_{\theta} f_t(\theta_{t-1})$ (Get gradients w.r.t. stochastic objective at timestep t)

$m_t \leftarrow \beta_1 \cdot m_{t-1} + (1 - \beta_1) \cdot g_t$ (Update biased first moment estimate)

$v_t \leftarrow \beta_2 \cdot v_{t-1} + (1 - \beta_2) \cdot g_t^2$ (Update biased second raw moment estimate)

$\hat{m}_t \leftarrow m_t / (1 - \beta_1^t)$ (Compute bias-corrected first moment estimate)

$\hat{v}_t \leftarrow v_t / (1 - \beta_2^t)$ (Compute bias-corrected second raw moment estimate)

$\theta_t \leftarrow \theta_{t-1} - \alpha \cdot \hat{m}_t / (\sqrt{\hat{v}_t} + \epsilon)$ (Update parameters) Also add a bias correction term

end while

return θ_t (Resulting parameters) for better conditioning in earlier iterations

Kingma, D. P., & Ba, J. (2014). Adam: A method for stochastic optimization. *arXiv preprint arXiv:1412.6980*.

Adaptive Learning Rate via ADAM

$$m_t := \alpha \cdot m_{t-1} + (1 - \alpha) \cdot \frac{\partial \mathcal{L}}{\partial w_{i,j}}(t)$$

$$r := \beta \cdot \text{MeanSquare}(w_{i,j}, t-1) + (1 - \beta) \left(\frac{\partial \mathcal{L}}{\partial w_{i,j}}(t) \right)^2$$

CLASS `torch.optim.Adam(params, lr=0.001, betas=(0.9, 0.999), eps=1e-08, weight_decay=0, amsgrad=False)`

[SOURCE] 


Implements Adam algorithm.

It has been proposed in [Adam: A Method for Stochastic Optimization](#).

The default settings for the "betas" work usually just fine

- Parameters:**
- **params** (*iterable*) – iterable of parameters to optimize or dicts defining parameter groups
 - **lr** (*float*, *optional*) – learning rate (default: 1e-3)
 - **betas** (*Tuple[float, float]*, *optional*) – coefficients used for computing running averages of gradient and its square (default: (0.9, 0.999))
 - **eps** (*float*, *optional*) – term added to the denominator to improve numerical stability (default: 1e-8)
 - **amsgrad** (*bool*, *optional*) – whether to use the AMSGrad variant of this algorithm from the paper [On the Convergence of Adam and Beyond](#) (default: False)

Source: <https://pytorch.org/docs/stable/optim.html>



<https://bl.ocks.org/EmilienDupont/aaf429be5705b219aaaf8d691e27ca87>

Improving Stochastic Gradient Descent

Lecture Overview

1. Applying "tricks" to HW3
2. Learning rate decay
3. Training with "momentum"
4. Adaptive learning rates
- 5. Using optimization algorithms in PyTorch**

Using Different Optimizers in PyTorch

Usage is the same as for vanilla SGD, which we used before,
you can find an overview at: <https://pytorch.org/docs/stable/optim.html>

```
optimizer = torch.optim.SGD(model.parameters(), lr=0.01, momentum=0.9)
optimizer = torch.optim.Adam(model.parameters(), lr=0.0001)
```

Using Different Optimizers in PyTorch

Usage is the same as for vanilla SGD, which we used before,
you can find an overview at: <https://pytorch.org/docs/stable/optim.html>

```
optimizer = torch.optim.SGD(model.parameters(), lr=0.01, momentum=0.9)
optimizer = torch.optim.Adam(model.parameters(), lr=0.0001)
```

Remember to save the optimizer state if you are using, e.g., Momentum or ADAM, and want to continue training later (see earlier slides on saving states of the learning rate schedulers).

Training Loss vs Generalization Error


Improving Generalization Performance by Switching from Adam to SGD

Nitish Shirish Keskar, Richard Socher


(Submitted on 20 Dec 2017)

Despite superior training outcomes, adaptive optimization methods such as Adam, Adagrad or RMSprop have been found to generalize poorly compared to Stochastic gradient descent (SGD). These methods tend to perform well in the initial portion of training but are outperformed by SGD at later stages of training. We investigate a hybrid strategy that begins training with an adaptive method and switches to SGD when appropriate. Concretely, we propose SWATS, a simple strategy which switches from Adam to SGD when a triggering condition is satisfied. The condition we propose relates to the projection of Adam steps on the gradient subspace. By design, the monitoring process for this condition adds very little overhead and does not increase the number of hyperparameters in the optimizer. We report experiments on several standard benchmarks such as: ResNet, SENet, DenseNet and PyramidNet for the CIFAR-10 and CIFAR-100 data sets, ResNet on the tiny-ImageNet data set and language modeling with recurrent networks on the PTB and WT2 data sets. The results show that our strategy is capable of closing the generalization gap between SGD and Adam on a majority of the tasks.

Training Loss vs Generalization Error



Kingma, D. P., & Ba, J. (2014). Adam: A method for stochastic optimization. *arXiv preprint arXiv:1412.6980*.



Keskar, N. S., & Socher, R. (2017). Improving generalization performance by switching from adam to sgd. *arXiv preprint arXiv:1712.07628*.

Reading Assignment

"An overview of gradient descent optimization algorithms" by Sebastian Ruder:

<http://ruder.io/optimizing-gradient-descent/index.html>

Stuff in the News!



```
import torch
import torchlayers

# torch.nn and torchlayers can be mixed easily
model = torch.nn.Sequential(
    torchlayers.Conv(64), # specify ONLY out_channels
    torch.nn.ReLU(), # use torch.nn wherever you wish
    torchlayers.BatchNorm(), # BatchNormNd inferred from input
    torchlayers.Conv(128), # Default kernel_size equal to 3
    torchlayers.ReLU(),
    torchlayers.Conv(256, kernel_size=11), # "same" padding as default
    torchlayers.GlobalMaxPool(), # Known from Keras
    torchlayers.Linear(10), # Output for 10 classes
)


print(model)
```

Above would give you model's summary like this (notice question marks for not yet inferred values):

```
Sequential(
  (0): Conv(in_channels=?, out_channels=64, kernel_size=3, stride=1, padding=same, dilation=1, groups=1, bias=True)
  (1): ReLU()
  (2): BatchNorm(num_features=?, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (3): Conv(in_channels=?, out_channels=128, kernel_size=3, stride=1, padding=same, dilation=1, groups=1, bias=True)
  (4): ReLU()
  (5): Conv(in_channels=?, out_channels=256, kernel_size=11, stride=1, padding=same, dilation=1, groups=1, bias=True)
  (6): GlobalMaxPool()
  (7): Linear(in_features=?, out_features=10, bias=True)
)
```

<https://github.com/szymonmaszke/torchlayers>

AnimeGAN



<https://github.com/TachibanaYoshino/AnimeGAN>

https://github.com/TachibanaYoshino/AnimeGAN/blob/master/doc/ISICA_2019_paper_7.pdf



<https://github.com/TachibanaYoshino/AnimeGAN>



Microsoft can filter out the sound of you eating potato chips on a conference call

<https://www.cnet.com/news/microsoft-can-filter-out-the-sound-of-you-eating-potato-chips-on-a-conference-call/>

Just in time for us all to be working from home because of the pandemic coronavirus, Microsoft's built new technology for its Teams software that can identify your voice and filter out any other sounds. That could be a leaf blower, a dog barking or even, yes, you rustling through a bag of potato chips.

A CONSTRUCTIVE PREDICTION OF THE GENERALIZATION ERROR ACROSS SCALES


Jonathan S. Rosenfeld¹ **Amir Rosenfeld²** **Yonatan Belinkov¹³** **Nir Shavit¹⁴⁵**
`{jonsr,belinkov,shanir}@csail.mit.edu` `amir@cse.yorku.ca`

¹ Massachusetts Institute of Technology ² York University ³ Harvard University
⁴ Neural Magic Inc ⁵ Tel Aviv University


ABSTRACT

The dependency of the generalization error of neural networks on model and dataset size is of critical importance both in practice and for understanding the theory of neural networks. Nevertheless, the functional form of this dependency remains elusive. In this work, we present a functional form which approximates well the generalization error in practice. Capitalizing on the successful concept of model scaling (e.g., width, depth), we are able to simultaneously construct such a form and specify the exact models which can attain it across model/data scales. Our construction follows insights obtained from observations conducted over a range of model/data scales, in various model types and datasets, in vision and language tasks. We show that the form both fits the observations well across scales, and provides accurate predictions from small- to large-scale models and data.

<https://openreview.net/pdf?id=ryenvpEKDr>



(a) Wiki103 error (cross entropy) landscape.



(b) CIFAR10 error (top1) landscape.

Figure 1: Error landscapes in log-log-log scale. Each point (blue dot) is the error resulting from training with a model/data configuration m, n . The surface is a linear interpolation between the points, which is then projected on the (m, ϵ) , (n, ϵ) and (m, n) planes. See Appendix C for details.



COVID-19 Open Research Dataset Challenge (CORD-19)

An AI challenge with AI2, CZI, MSR, Georgetown, NIH & The White House



Allen Institute For AI and 8 collaborators • updated 3 days ago (Version 5)

Data

Tasks (10)

Kernels (509)

Discussion (244)

Activity

Metadata

Download (4 GB)

New Notebook

⋮

Dataset Description

In response to the COVID-19 pandemic, the White House and a coalition of leading research groups have prepared the COVID-19 Open Research Dataset (CORD-19). CORD-19 is a resource of over 45,000 scholarly articles, including over 33,000 with full text, about COVID-19, SARS-CoV-2, and related coronaviruses. This freely available dataset is provided to the global research community to apply recent advances in natural language processing and other AI techniques to generate new insights in support of the ongoing fight against this infectious disease. There is a growing urgency for these approaches because of the rapid acceleration in new coronavirus literature, making it difficult for the medical research community to keep up.

<https://www.kaggle.com/allen-institute-for-ai/CORD-19-research-challenge>